

Data Structures and Algorithms

Lecture 8: Sorting

Outline of Today's Lecture

- Internal sorting
 - Three basic sorting algorithms ← $\Theta(n^2)$
 - Insertion / bubble / selection sorts
 - One medium sorting algorithms ← $\Theta(n^{1.5})$
 - Shell sort
 - Three fast sorting algorithms ← $\Theta(n \log n)$
 - Merge / quick / heap sorts
 - Two special cases ← $\Theta(n)$
 - Bin / radix sorts
- External sorting

Sorting

- **Motivation:** Suppose the record of student consists of student name, ID, course name, score, we sort n students by their **scores**.
- Given a set of records r_1, r_2, \dots, r_n with key values k_1, k_2, \dots, k_n , the **Sorting Problem** is to arrange the records in non-decreasing order by their keys.
- Measures of algorithm cost:
 - Comparisons and Swaps are two main operations in sorting.

Sorting terminology

- Input is a set of records stored *in an array*.
- A sorting algorithm is **stable**, if it does not change the relative ordering of records with identical key values.
- **Internal sorting vs. External sorting**
 - In interval sorting, all records can be loaded into a computer memory
 - External sorting, there are **too many records** to be sorted → **cannot** be loaded into the memory

Internal Sorting Algorithms

- Three $\Theta(n^2)$ sorting algorithms
 - Insertion / bubble / selection sorts
- Shell sort -- $O(n^{1.5})$ in average case
- Three quick sorting algorithms-- $\Theta(n \log n)$
 - Merge / quick / heap sorts
- Two $\Theta(n)$ sorting algorithms for special cases of record keys
- Lower bounds for sorting

Insertion Sort (1)

- Assume you have sorted the first i (e.g., $i=2$) numbers, consider the $(i+1)$ th number 36, **insert** the number **in order** so that the first $i+1$ numbers are sorted.

Before insert: [27 53] 36 15 69 42

After insert : [27 36 53] 15 69 42

Insertion Sort (2)

- Traverse i from 1 to $n-1$, do the insertion

$i=1:$	[53]	27	36	15	69	42
$i=2:$	[27	53]	36	15	69	42
$i=3:$	[27	36	53]	15	69	42
$i=4:$	[15	27	36	53]	69	42
$i=5:$	[15	27	36	53	69]	42
	[15	27	36	42	53	69]

Insertion Sort (3)

```
template <class E>
void insertSort(E A[], int n) {
    for (int i=1; i<n; i++)
        for (int j=i; j>0 && A[j] < A[j-1]; j--)
            swap(A, j, j-1);
}
```

Best Case Analysis of Insertion Sort

- The best case occurs when the initial list of numbers are already sorted
- Best Case: 0 swap, $n - 1$ comparisons

```
15  27  36  42  53  69
```

Worst Case Analysis of Insertion Sort

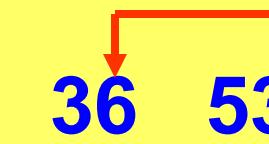
- The worst case occurs when the initial list of numbers are **reversely** sorted
- At i -th iteration, performs i comparisons and swaps
- Total: $\Sigma i = n^2/2$ swaps and comparisons

69 53 42 36 27 15

Average Case Analysis of Insertion Sort

- At i -th iteration, performs $i/2$ comparisons and swaps **on average**
- Total: $\Sigma i/2 = n^2/4$ swaps and comparisons

Before insert:	[27	53]	36	15	69	42
After insert :	[27	36	53]	15	69	42



Insertion Sort

- Best Case: 0 swap, $n - 1$ comparisons
- Worst Case: $n^2/2$ swaps and comparisons
- Average Case: $n^2/4$ swaps and comparisons
- Insertion Sort is **suitable** for the cases where the **records** in the input array are **almost sorted**, e.g.,
 - Many records are already been sorted initially, but some **a few new records** are added

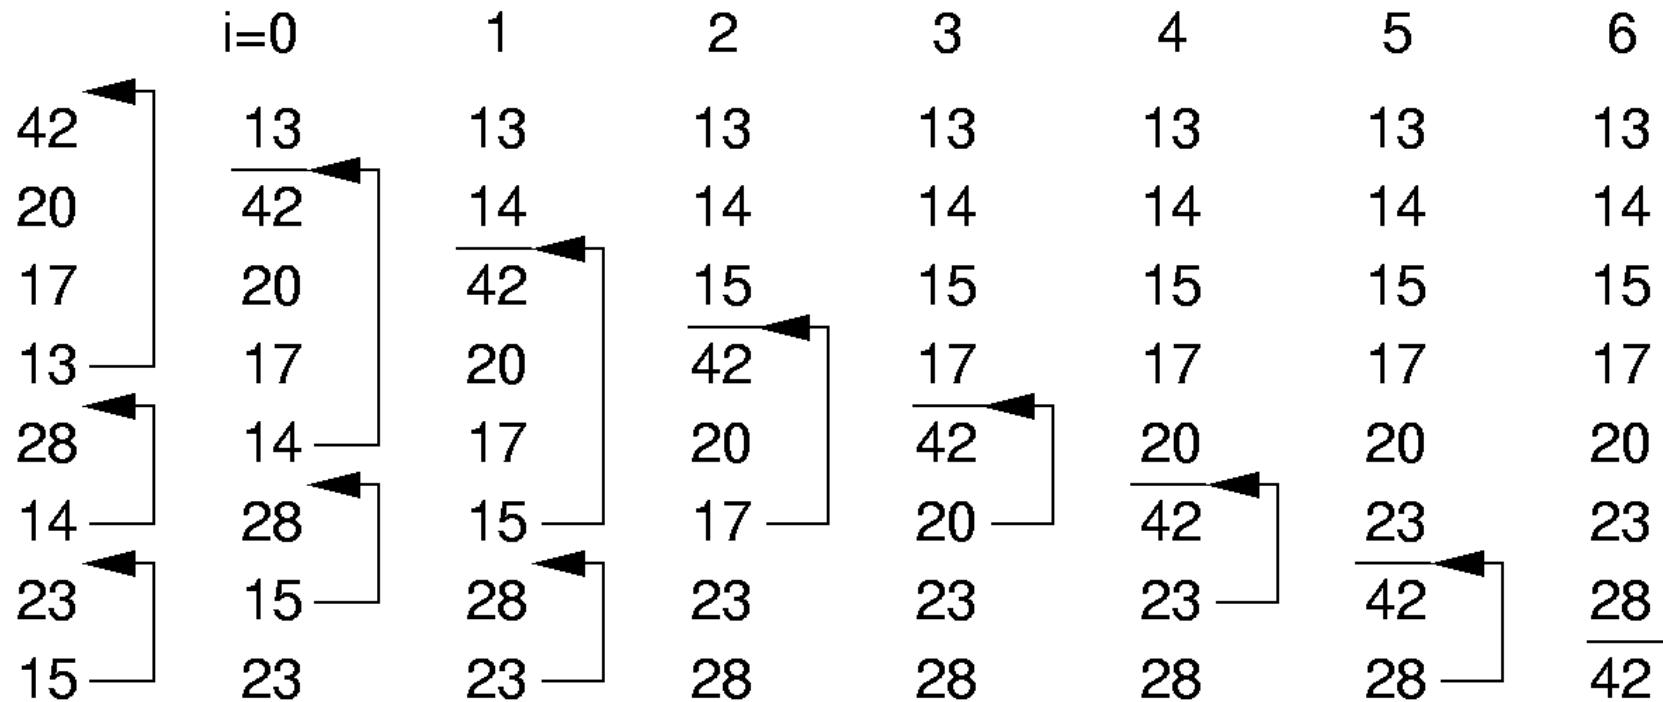
Bubble Sort (1)

- Scan from the bottom to the top, compare each adjacent values $K[j-1]$ and $K[j]$, swap them if the $K[j] < K[j-1]$. After the scan, the *smallest* value is at the top (**bubble up**)
- Do the 2nd scan from the bottom to the top-2

- ...

	i=0	i=0	1
42	13	13	13
20	42	42	14
17	20	20	42
13	17	17	20
28	14	14	17
14	28	28	15
23	15	15	28
15	23	23	23

Bubble Sort (2)



Bubble Sort (3)

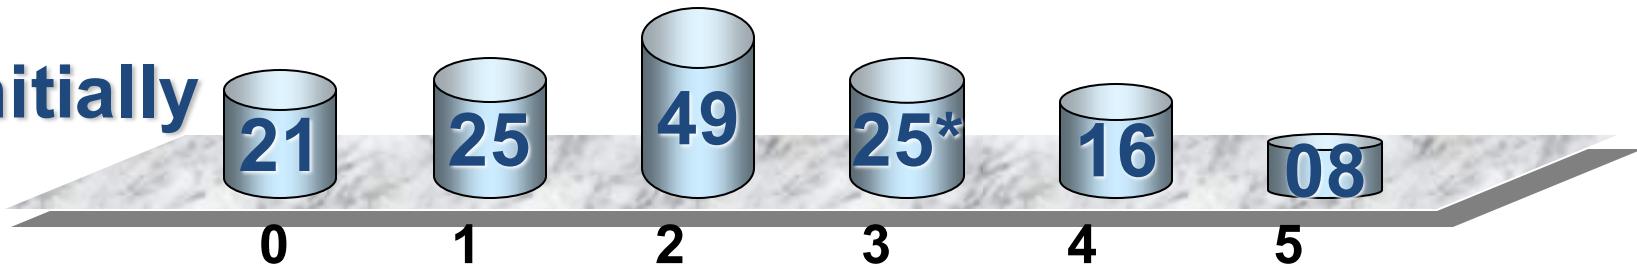
```
template <class E>
void bubbleSort(E A[], int n) {
    for (int i=0; i<n-1; i++)
        for (int j=n-1; j>i; j--)
            if (A[j] < A[j-1])
                swap(A, j, j-1);
}
```

- Best Case: 0 swaps, $n^2/2$ comparisons
- Worst Case: $n^2/2$ swaps and comparisons
- Average Case: $n^2/4$ swaps and $n^2/2$ comparisons

Selection Sort

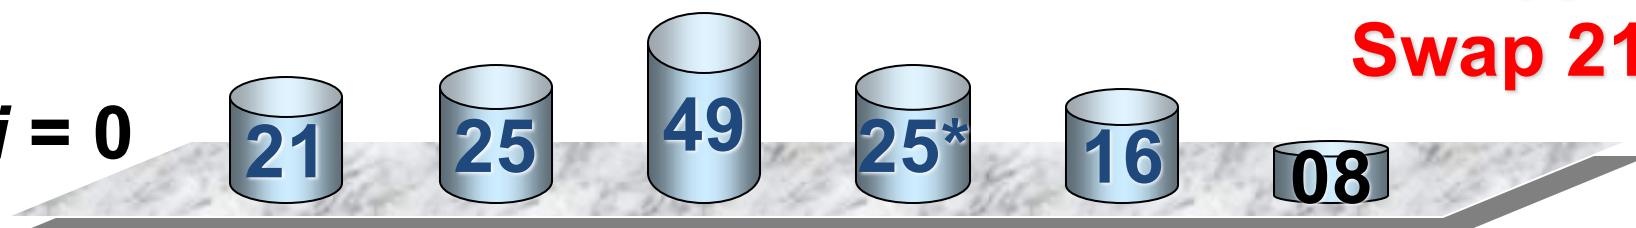
- Basic idea:
- First, **select** the **smallest** value, store it at the **first** location in the array
- Select the **2nd smallest** value, store it at the **2nd location** in the array
- ...
- The array is sorted after ***n*** iterations

initially



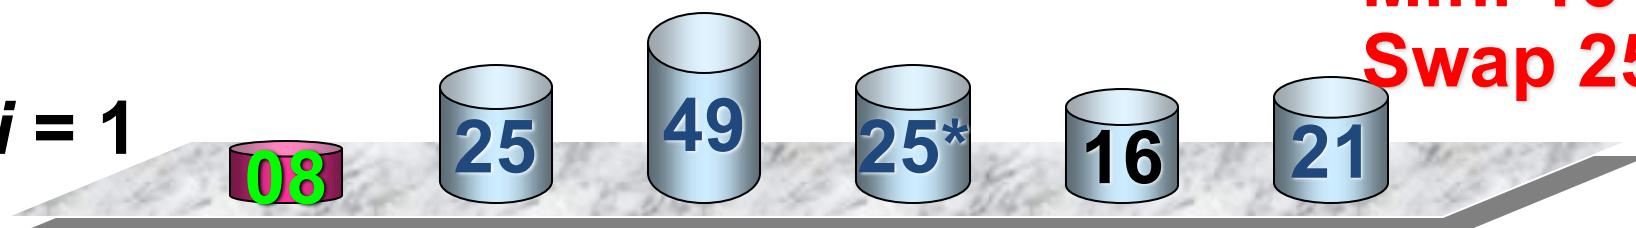
Min:08
Swap 21,08

$i = 0$



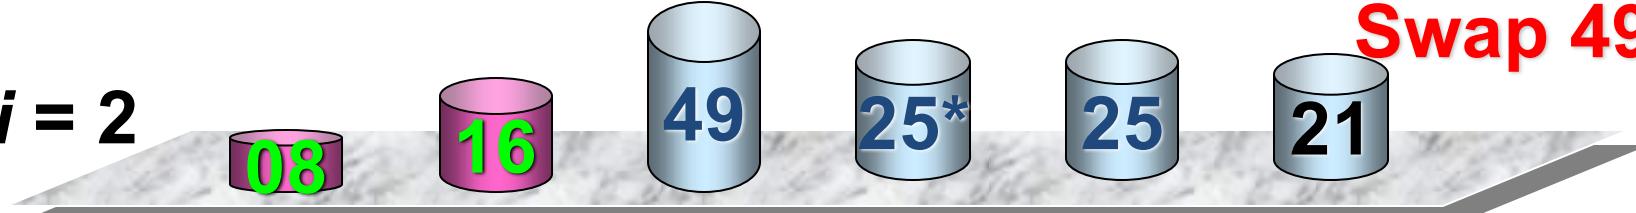
Min: 16
Swap 25,16

$i = 1$

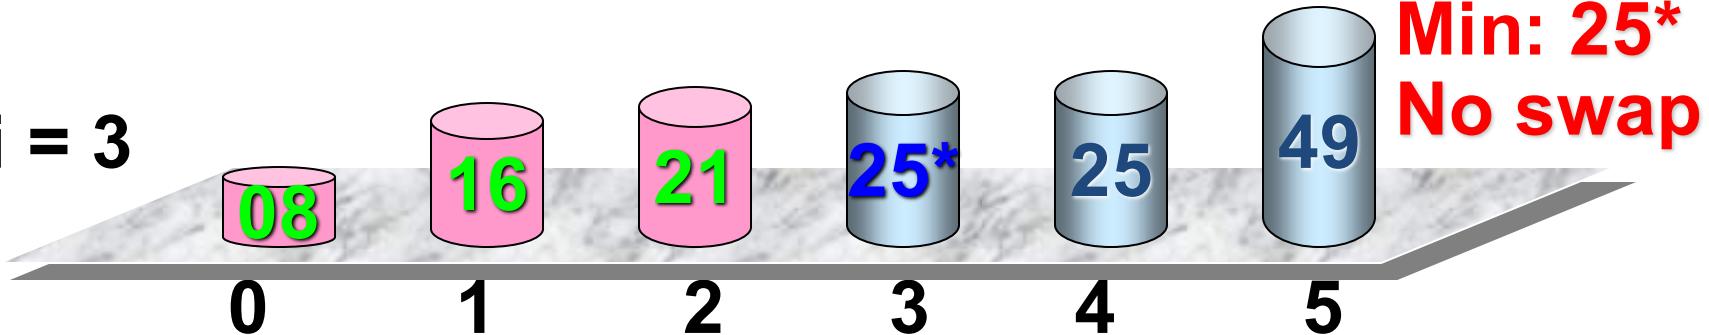


Min: 21
Swap 49,21

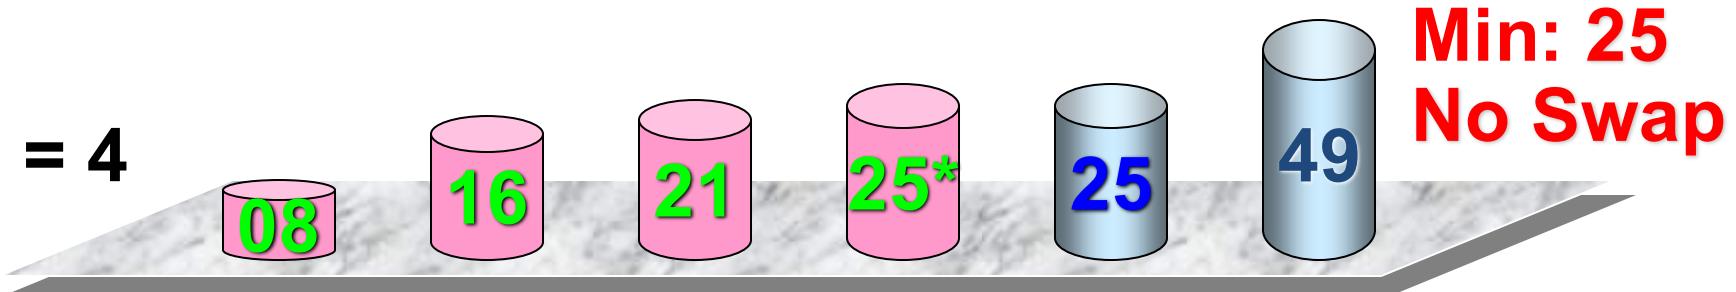
$i = 2$



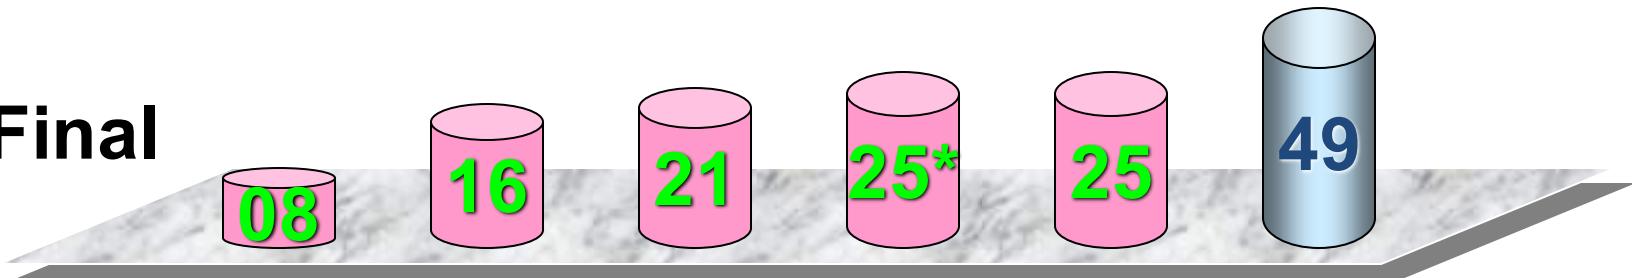
$i = 3$



$i = 4$



Final



Selection Sort (2)

```
template <class E>
void selectionSort(E A[], int n) {
    for (int i=0; i<n-1; i++) {
        int lowindex = i; // Remember its index
        for (int j=n-1; j>i; j--) // Find least
            if (A[j] < A[lowindex])
                lowindex = j; // Put it in place
        swap(A, i, lowindex);
    }
}
```

- Best case: $n-1$ swaps, $n^2/2$ comparisons.
- Worst case: $n - 1$ swaps and $n^2/2$ comparisons.
- Average case: $n-1$ swaps and $n^2/2$ comparisons.

Summary of three $\Theta(n^2)$ sorting algorithms

	Insertion	Bubble	Selection
--	-----------	--------	-----------

Comparisons:

Best Case	$\Theta(n)$	$\Theta(n^2)$	$\Theta(n^2)$
Average Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n^2)$
Worst Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n^2)$

Swaps:

Best Case	0	0	$\Theta(n)$
Average Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n)$
Worst Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n)$

Running time comparisons ($n=100k$)

■ Random input

100000 numbers to be sorted:

Sort with InsertionSort, Time consumed: 1.438 (seconds)

Sort with bubble sort, Time consumed: 17.360 (seconds)

Sort with SelectionSort, Time consumed: 3.813 (seconds)

■ The input is already sorted

100000 numbers to be sorted:

Sort with InsertionSort, Time consumed: 0.000 (seconds)

Sort with bubble sort, Time consumed: 3.766 (seconds)

Sort with SelectionSort, Time consumed: 3.905 (seconds)

■ The input is reversely sorted

100000 numbers to be sorted:

Sort with InsertionSort, Time consumed: 2.984 (seconds)

Sort with bubble sort, Time consumed: 9.692 (seconds)

Sort with SelectionSort, Time consumed: 3.872 (seconds)

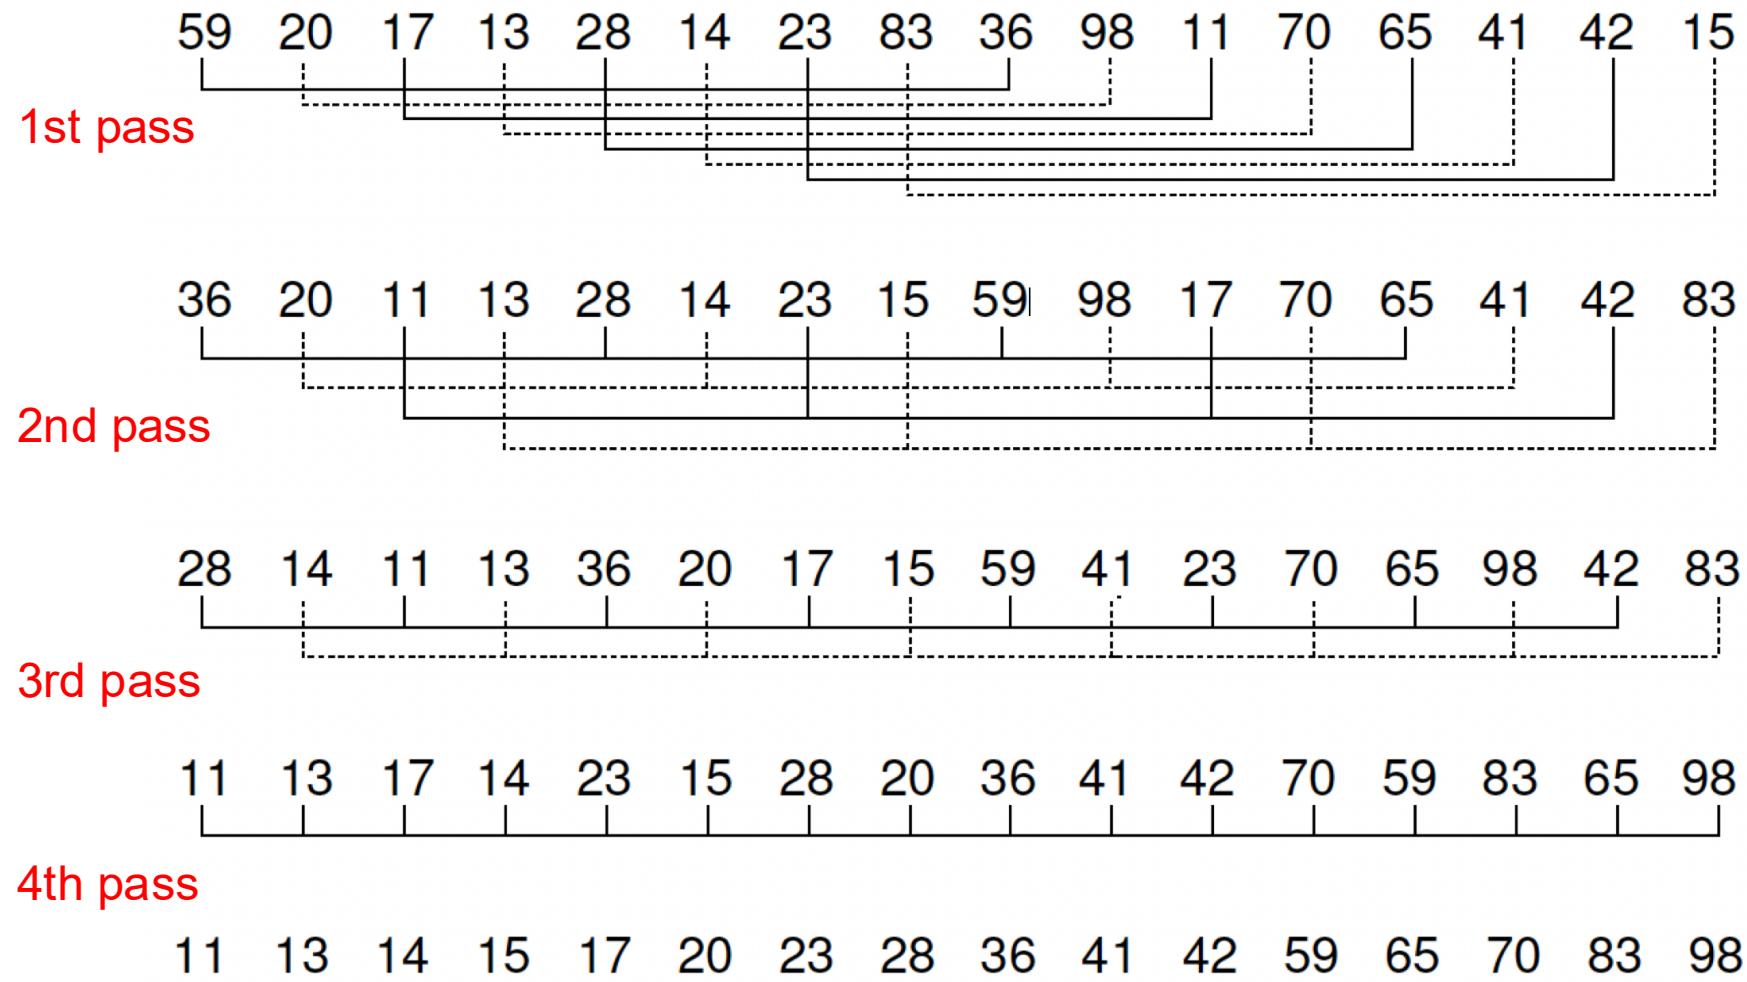
Shell sort

- **Shellsort**, named after its inventor, D.L. Shell.
 - Sometimes called the **diminishing increment** sort.
 - $O(n^{1.5})$ on average-case
- Its **strategy** is to make the list "mostly sorted" so that a final Insertion Sort can finish the job.
- Main steps:
 - Break the list into sublists
 - Sort them
 - Then, recombine the sublists

Shell sort process

- During each iteration/pass, Shellsort breaks the list into **disjoint sublists** so that each element in a sublist is a fixed number of positions apart. e.g.,
 - Let us assume for convenience that n , the number of values to be sorted, is *a power of two*.
 - Shellsort will begin by breaking the list into $n/2$ sublists of 2 elements each, where the array index of the 2 elements in each sublist differs by $n/2$.

Shell sort - an example (1)



Shell sort - an example (2)

- Some choices for increments would make Shellsort run more efficiently.
 - In particular, the choice of increments described above ($2^k, 2^{k-1}, \dots, 2, 1$) turns out to be relatively inefficient.
 - A better choice is the following series based on devision by three: (... , 121, 40, 13, 4, 1).

Shellsort Implementation

```
// Modified version of Insertion Sort for varying increments
template <typename E, typename Comp>
void inssort2(E A[], int n, int incr) {
    for (int i=incr; i<n; i+=incr)
        for (int j=i; (j>=incr) &&
            (Comp::prior(A[j], A[j-incr]))); j-=incr)
            swap(A, j, j-incr);
}

template <typename E, typename Comp>
void shellsort(E A[], int n) { // Shellsort
    for (int i=n/2; i>2; i/=2) // For each increment
        for (int j=0; j<i; j++) // Sort each sublist
            inssort2<E,Comp>(&A[j], n-j, i);
    inssort2<E,Comp>(A, n, 1);
}
```

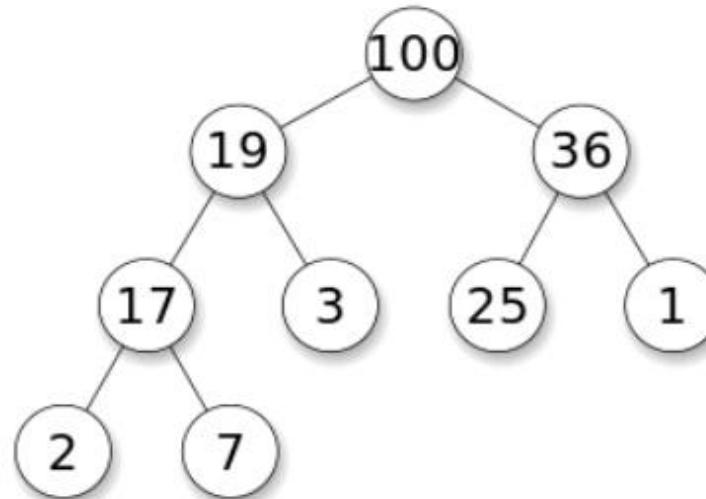
Three fast sorting algorithms

- Heap sort
 - $\Theta(n \log n)$ for the worst, best, average cases
- Merge sort
 - $\Theta(n \log n)$ for the worst, best, average cases
- Quick sort
 - $\Theta(n \log n)$ for the best and average cases
 - $\Theta(n^2)$ for the worst case

Heap – a special binary tree (Ch. II.5)

Heap: Complete binary tree with the heap property:

- Max-heap: each value in a node is no less than its children values
- The values in the tree are partially ordered.
 - The left child may less or greater than its right child

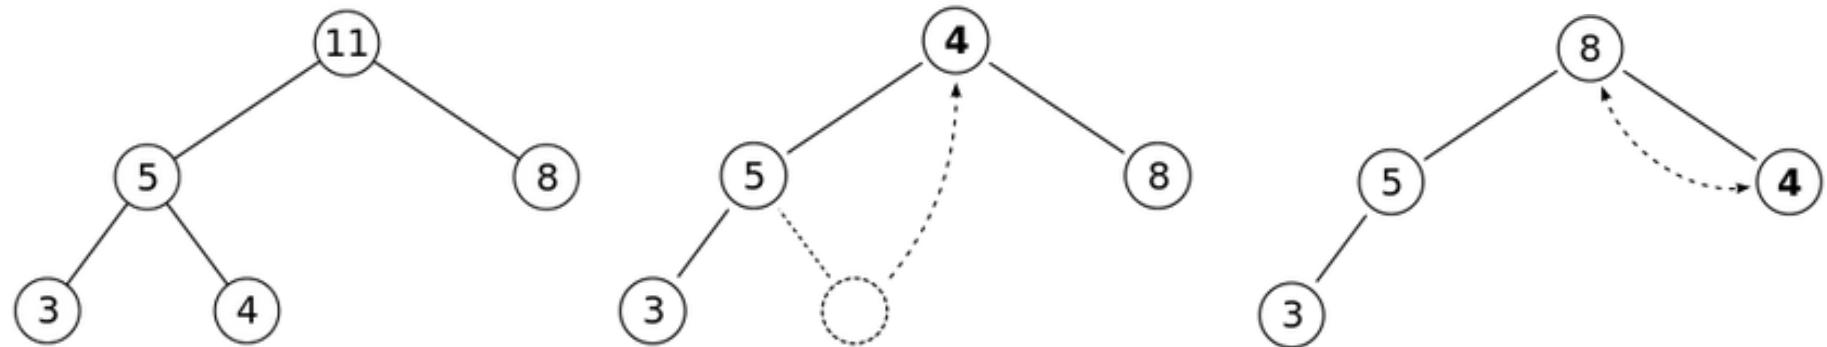


Heap Sort

- Given an array, build a max-heap
- Remove the **maximum** number from the heap
- Remove the **next maximum** number
- ...
- Continue until no numbers are left in the heap

Heap -- removeMax

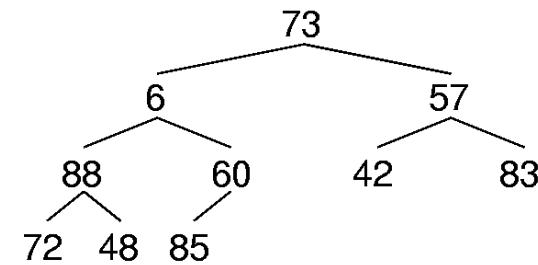
- Replace the **root** of the heap with the **last** element on the last level
- Compare the **new root** with its children (**shift down operation**)
 - if the new root is **larger** than its children, stop.
 - If not, swap the element with its **largest children**, and return to the previous step
 - **Worst time complexity $\Theta(\log n)$**



HeapSort Example (1)

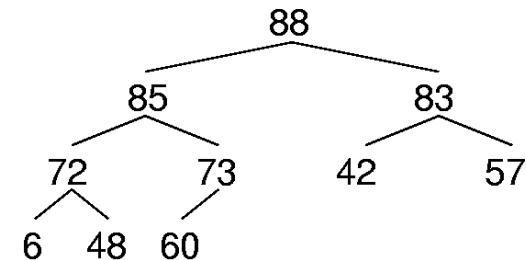
Original Numbers

73	6	57	88	60	42	83	72	48	85
----	---	----	----	----	----	----	----	----	----



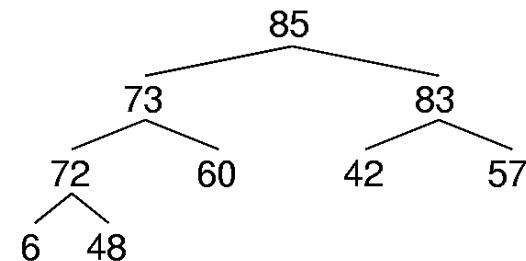
Build Heap

88	85	83	72	73	42	57	6	48	60
----	----	----	----	----	----	----	---	----	----



Remove 88

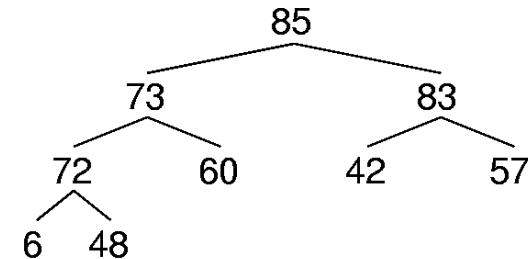
85	73	83	72	60	42	57	6	48	88
----	----	----	----	----	----	----	---	----	----



HeapSort Example (2)

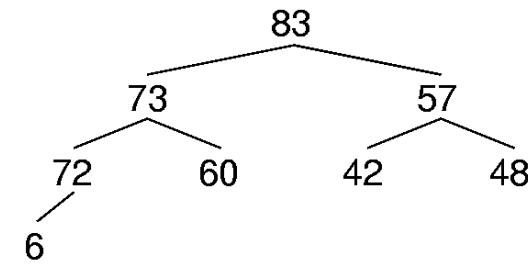
Remove 88

85	73	83	72	60	42	57	6	48	88
----	----	----	----	----	----	----	---	----	----



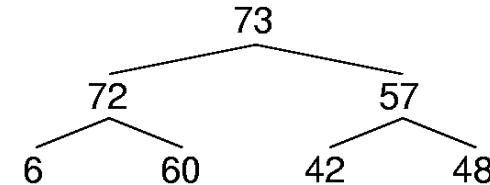
Remove 85

83	73	57	72	60	42	48	6	85	88
----	----	----	----	----	----	----	---	----	----



Remove 83

73	72	57	6	60	42	48	83	85	88
----	----	----	---	----	----	----	----	----	----



Heapsort

```
template <class E>
void heapSort(E A[], int n) { // Heapsort
    E mval;
    maxheap<E> H(A, n, n);
    for (int i=0; i<n; i++) // Now sort
        H.removemax(mval); // Put max at end
}
```

Analysis of Heap Sort

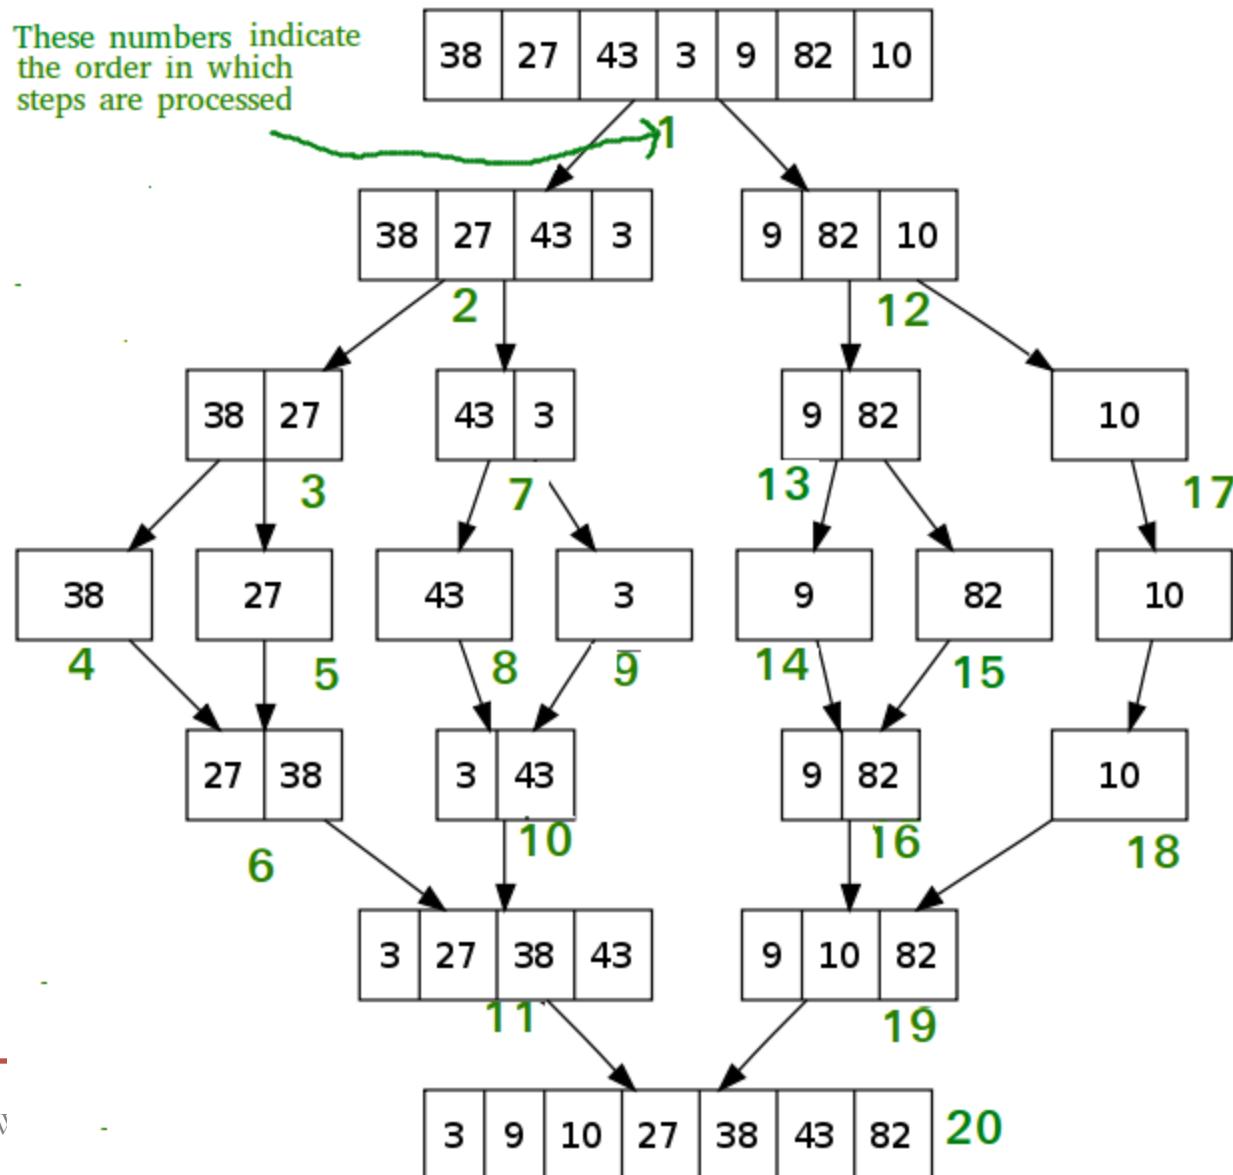
- Build a heap takes time $\Theta(n)$
- Remove the maximum value takes $\Theta(\log n)$,
as heap is a complete tree
- Total time is $\Theta(n) + n \Theta(\log n) = \Theta(n \log n)$

Merge Sort

- Basic idea: **divide and conquer**
 1. Given a list of numbers to be sorted
 2. **Split** the list into **two sub-lists** with the identical length
 3. Recursively sort the sub-lists, respectively
 4. **Merge** the two sorted sub-lists

Merge Sort

These numbers indicate
the order in which
steps are processed



Merge sort with an **array-based** list (1)

- An array $A[\text{left}, \dots, \text{right}]$, with the index range:
 $\text{left} \dots \text{right}$
- How to split ?
- Let $\text{mid} = (\text{left} + \text{right})/2$
- Left sub-list = $A[\text{left}, \dots, \text{mid}]$
- Right sub-list= $A[\text{mid+1}, \dots, \text{right}]$

Merge Sort with an array-based list (2)

- How to merge two sorted sub-lists $A[\text{left}, \dots, \text{mid}]$, $A[\text{mid}+1, \dots, \text{right}]$?



- An extra array $\text{temp}[\text{left} \dots, \text{right}]$ is needed



- Step 1: move the smallest value of the first numbers of the two-sublists to array temp
 - If one sub-list is exhausted, just move the first number of the other sublist
- Continue until no numbers are left
- Copy back: $A[\text{left}, \dots, \text{right}] = \text{temp}[\text{left}, \dots, \text{right}]$

Merge Sort Implementation

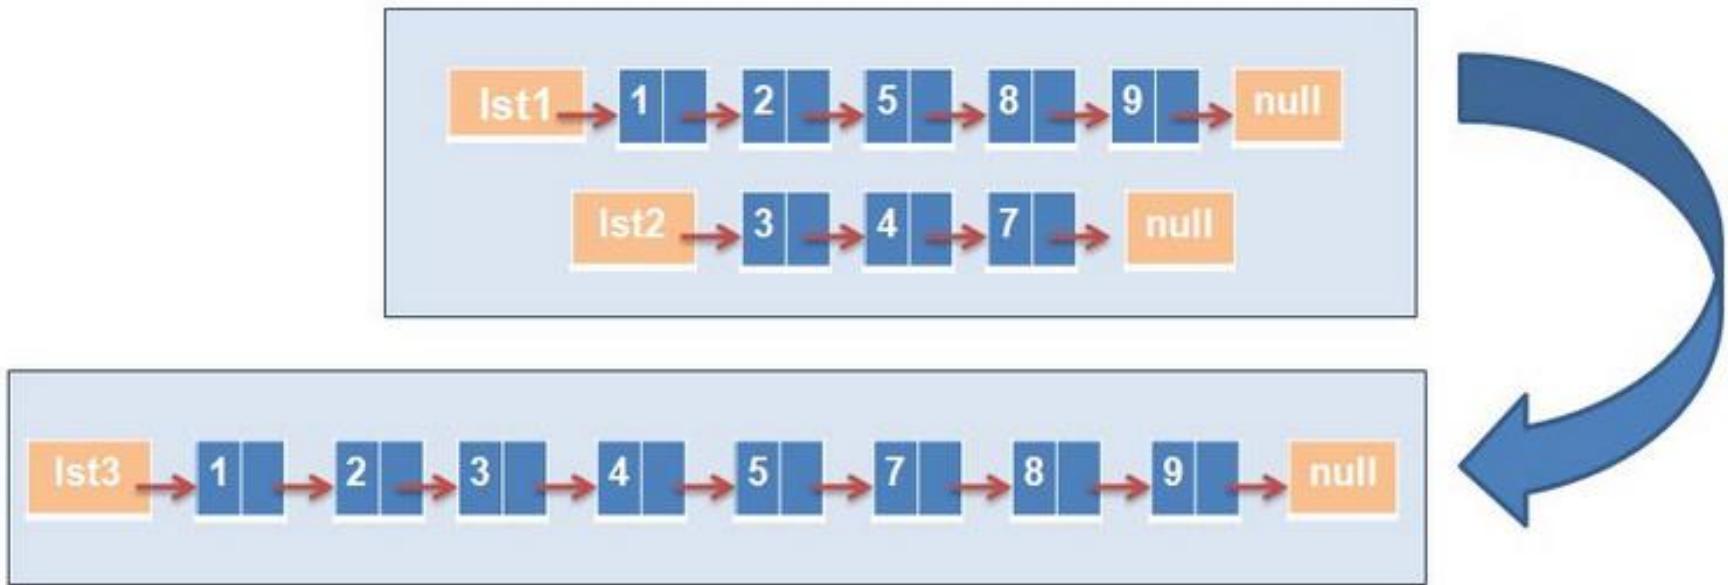
```
template <class E>
void mergeSort(E A[], E temp[],
                int left, int right) {
    if (left == right) return;
    int mid = (left+right)/2;
    mergesort<E>(A, temp, left, mid);
    mergesort<E>(A, temp, mid+1, right);
    //merge two sorted sublists
    int i1 = left; int i2 = mid + 1;
    for (int curr=left; curr<=right; curr++) {
        if (i1 == mid+1)           // Left exhausted
            temp[curr] = A[i2++];
        else if (i2 > right)      // Right exhausted
            temp[curr] = A[i1++];
        else if (A[i1] < A[i2])
            temp[curr] = A[i1++];
        else temp[curr] = A[i2++];
    }
    for (int i=left; i<=right; i++) // Copy back
        A[i] = temp[i];
}
```

Merge Sort based on a linked list (1)

- How to split ?
- Given a singly linked list of numbers
- Need to scan half of numbers in the list

How to merge two sorted sub-lists ?

- Similar to the array-based version
- But no extra memory is needed



Time complexity of Merge Sort

- Let $T(n)$ be the running time for n numbers
- Split: $\Theta(1)$ for the array-based list
- Recursively sorting two sub-lists: $2 * T(n/2)$
- Merge: $\Theta(n)$
- $T(n) = 2 T(n/2) + \Theta(n)$
- Expand the recurrence relationship, we have:
- $T(n) = \Theta(n \log n)$

Quick Sort

- Given an array of numbers $A[\text{left}, \dots, \text{right}]$
- Pick a value in the array as a **pivot**

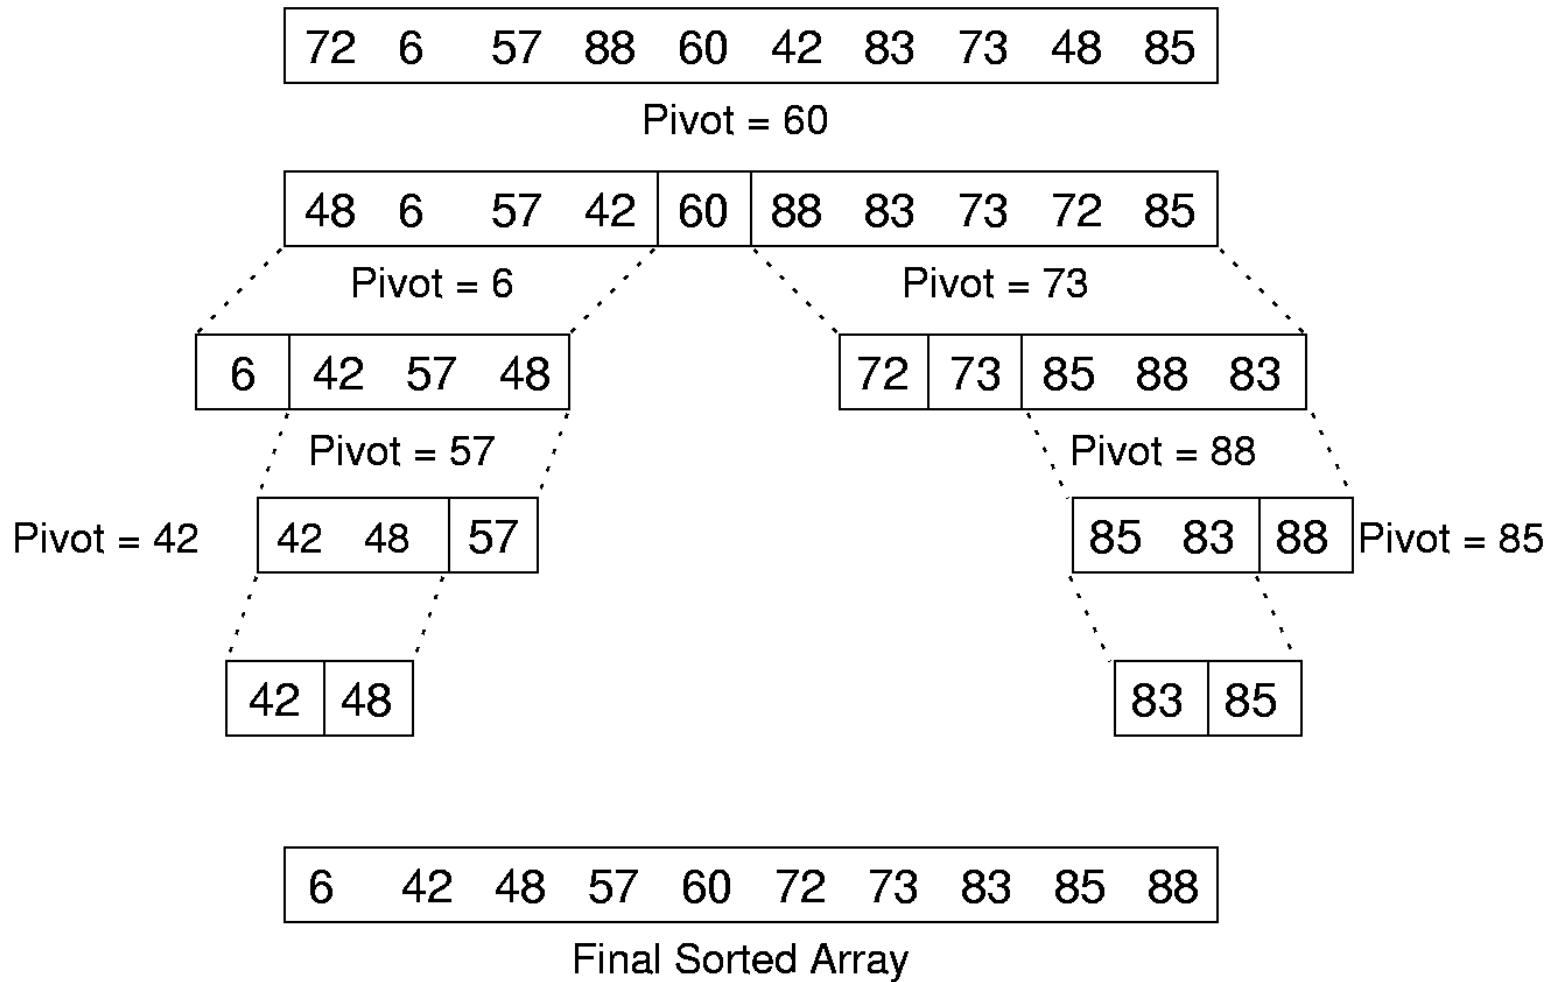


- Partition the array into **three parts**



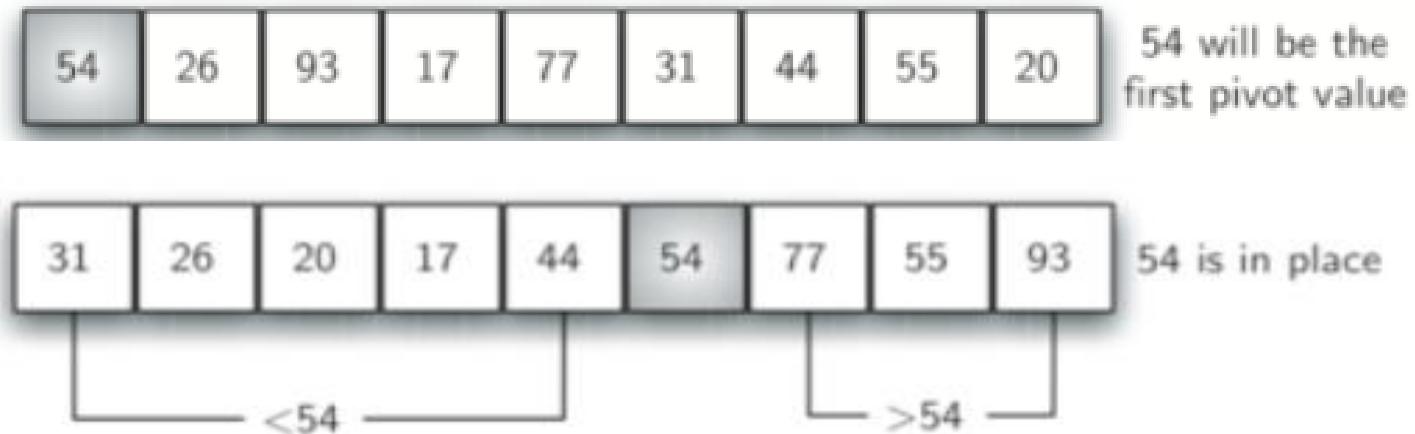
1. The numbers in the **left** part are $<$ pivot 54
 2. The pivot **itself** in place
 3. The numbers in the **right** part are \geq pivot 54
- **Recursively** sort the left and right parts

QuickSort Example



Two key problems in Quick Sort

- How to choose the pivot, such that the left and right parts are roughly **balanced** ?
 - The number of records in the left part is **more or less** the number in the right part
- How to **efficiently** partition an array by the pivot?



Solutions to the choice of a pivot

1. Traditionally, choose the **first** or the **last** number in the array
 - This is bad if the given array are already (or **nearly sorted**, or **reversely sorted**, one part has **0** number, the other part has **(n-1)** numbers
2. Choose the middle number
 - $\text{mid} = (\text{left}+\text{right})/2$; **pivot = A[mid];**
 - a better choice
3. median of three
 - Choose the pivot as the **median** of the first, middle and last numbers
 - Much better

Solutions to the choice of a pivot

4. Randomly choose a number as the pivot

- ❑ It is unlikely that a randomly chosen number is the smallest or the largest ones
- ❑ Can combine with the 3rd solution, i.e., **randomly choose three numbers**, and select the **median of the three** as the pivot

Partition an array by a given pivot

- Assume that the pivot is the median of the first, middle, and last numbers

pivot =



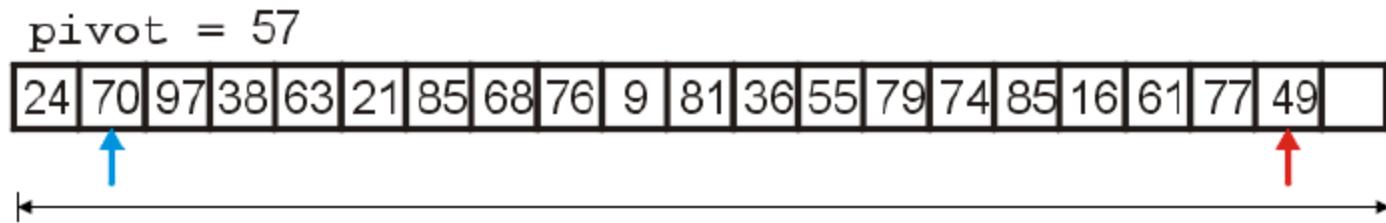
- First swap the pivot with the last number

pivot = 57

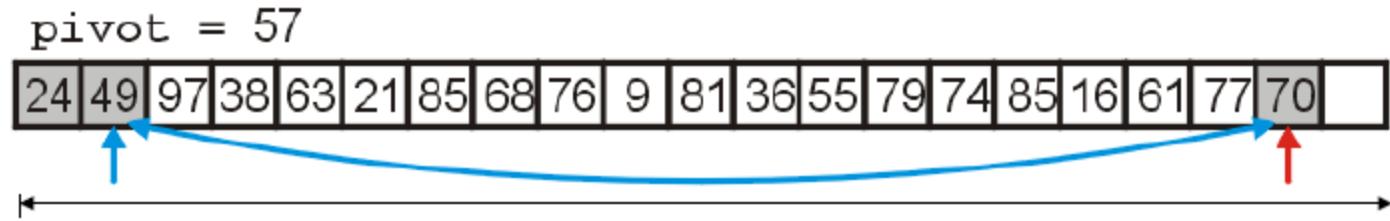


Partition an array by a given pivot

1. Start from the 1st location, search **forward** until we find a value \geq pivot, e.g., 70 > 57
2. Start from the 2nd last location, search **backward** until we find a value $<$ pivot, e.g., 49 < 57



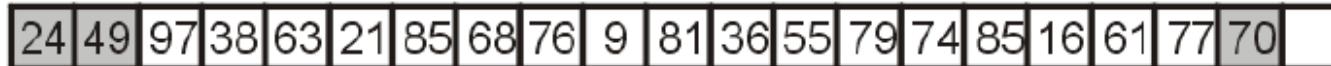
3. 70 and 49 are out of order, swap them



- We continue step1—step 3, see the next slides

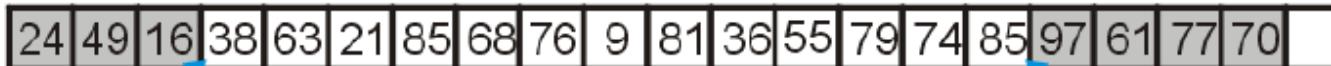
4. search forward until we find $97 \geq 57$
5. search backward until we find $16 < 57$

pivot = 57



6. Swap 97 and 16

pivot = 57



7. search forward until we find $63 \geq 57$

8. search backward until we find $55 < 57$

pivot = 57



9. Swap 63 and 55

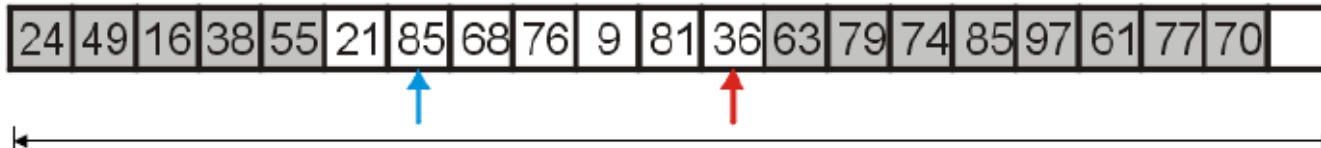
pivot = 57



10. search forward until we find $85 \geq 57$

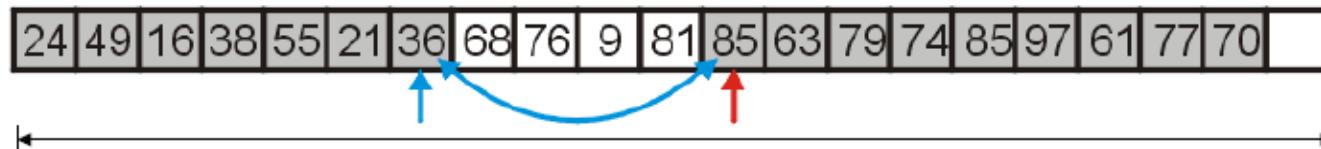
11. search backward until we find $36 < 57$

pivot = 57



12. Swap 85 and 36

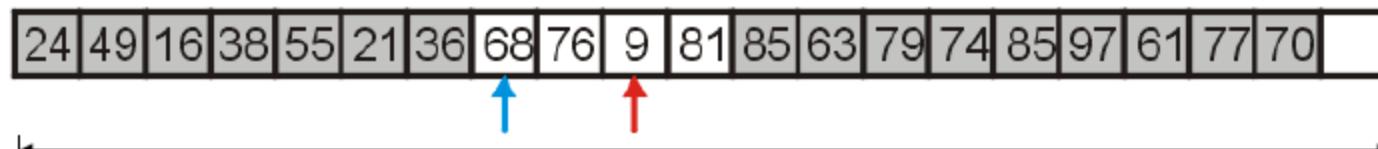
pivot = 57



13. search forward until we find $68 \geq 57$

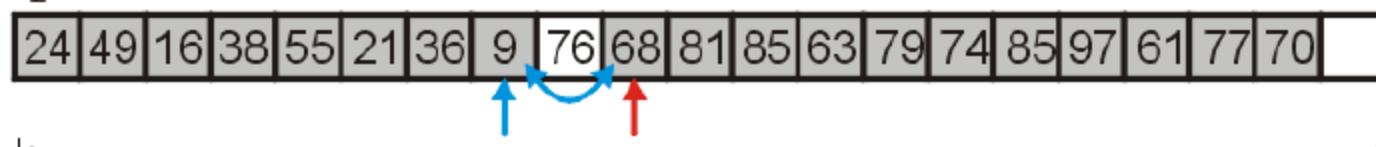
14. We search backward until we find $9 < 57$

pivot = 57



15. Swap 68 and 9

pivot = 57

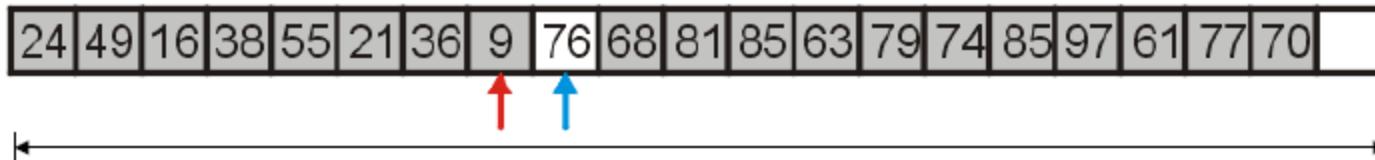


16. search forward until we find $76 \geq 57$

17. search backward until we find $9 < 57$

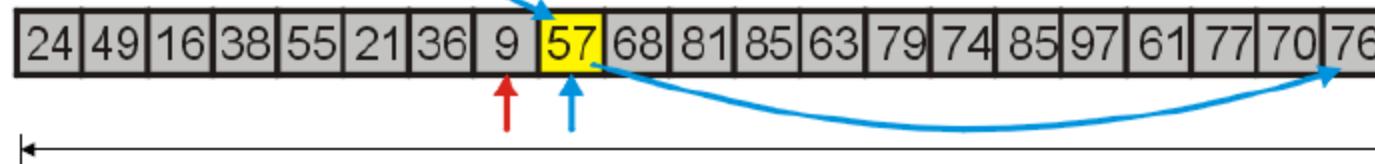
- The indices are out of order, stop

`pivot = 57`



- Move the first value larger than pivot, i.e., **76**, to the last location of the array
- Fill the empty location with the pivot **57**
- The pivot is in the correct location

`pivot = 57`



Another example of the partition ([animation](#))

Unsorted Array



Time complexity of Quick Sort

- Finding the pivot takes time $\Theta(1)$
- Partitioning an array takes time $\Theta(n)$
- Worst case time complexity
 - For each partition, one part has 0 number, the other has $n-1$ numbers
 - $T(n) = \Theta(n) + T(n-1)$
 - $T(n) = \Theta(n^2)$

Time complexity of Quick Sort

■ Best case analysis

- The best case occurs if the left and right parts are balanced, each has about $n/2$ numbers
- $T(n) = \Theta(n) + 2T(n/2)$
- $T(n) = \Theta(n \log n)$

Average time complexity of quick sort

- Consider all cases of the lengths of the two parts
 - Left: 0 number, right: $n-1$ numbers
 - Left: 1 number, right: $n-2$ numbers
 - Left: 2 number, right: $n-3$ numbers
 - ...
 - Left: $n-1$ number, right: 0 numbers
- Assume the probabilities of different cases are equal, i.e., $1/n$, we have

$$T(n) = cn + \frac{1}{n} \sum_{k=0}^{n-1} [T(k) + T(n-1-k)], \quad T(0) = T(1) = c.$$

$$T(n) = cn + \frac{1}{n} \sum_{k=0}^{n-1} [T(k) + T(n-1-k)] = cn + \frac{2}{n} \sum_{k=0}^{n-1} T(k)$$

|

$$nT(n) = cn^2 + 2 \sum_{k=0}^{n-1} T(k)$$

$$(n-1)T(n-1) = c(n-1)^2 + 2 \sum_{k=0}^{n-2} T(k)$$

$$nT(n) - (n-1)T(n-1) = c(2n-1) + 2T(n-1)$$

$$nT(n) = (n+1)T(n-1) + c(2n-1)$$

$$\begin{aligned} \frac{T(n)}{n+1} &= \frac{T(n-1)}{n} + \frac{c(2n-1)}{n(n+1)} \\ &\leq \frac{T(n-1)}{n} + \frac{c2n}{n(n+1)} \\ &= \frac{T(n-1)}{n} + \frac{2c}{n+1} \\ &= \frac{T(n-2)}{n-1} + \frac{2c}{n} + \frac{2c}{n+1} \end{aligned}$$

⋮

$$= \frac{T(1)}{2} + \sum_{i=2}^n \frac{2c}{i+1}$$

$$T(n) = \Theta(n \log n)$$

$$@ \text{CS311, H} \leq 2c \sum_{i=1}^{n-1} \frac{1}{i} \approx 2c \int_1^n \frac{1}{x} dx = 2c \ln n$$

Running time comparisons ($n=100k$)

■ Random input

```
100000 numbers to be sorted:  
Sort with InsertionSort, Time consumed: 1.438 (seconds)  
Sort with bubble sort, Time consumed: 17.360 (seconds)  
Sort with SelectionSort, Time consumed: 3.813 (seconds)  
Sort with shellSort, Time consumed: 0.016 (seconds)  
Sort with heapSort, Time consumed: 0.000 (seconds)  
Sort with mergeSort, Time consumed: 0.015 (seconds)  
Sort with quickSort, Time consumed: 0.016 (seconds)
```

■ The input is already sorted

```
100000 numbers to be sorted:  
Sort with InsertionSort, Time consumed: 0.000 (seconds)  
Sort with bubble sort, Time consumed: 3.766 (seconds)  
Sort with SelectionSort, Time consumed: 3.905 (seconds)  
Sort with shellSort, Time consumed: 0.002 (seconds)  
Sort with heapSort, Time consumed: 0.005 (seconds)  
Sort with mergeSort, Time consumed: 0.004 (seconds)  
Sort with quickSort, Time consumed: 0.000 (seconds)
```

■ The input is reversely sorted

```
100000 numbers to be sorted:  
Sort with InsertionSort, Time consumed: 2.984 (seconds)  
Sort with bubble sort, Time consumed: 9.692 (seconds)  
Sort with SelectionSort, Time consumed: 3.872 (seconds)  
Sort with shellSort, Time consumed: 0.004 (seconds)  
Sort with heapSort, Time consumed: 0.005 (seconds)  
Sort with mergeSort, Time consumed: 0.005 (seconds)  
Sort with quickSort, Time consumed: 0.001 (seconds)
```

Running time comparisons ($n=3M$)

- Random input

3000000 numbers to be sorted:	
Sort with shellSort, Time consumed:	1.087 (seconds)
Sort with heapSort, Time consumed:	0.860 (seconds)
Sort with mergeSort, Time consumed:	0.421 (seconds)
Sort with quickSort, Time consumed:	0.334 (seconds)

- The input is already sorted

3000000 numbers to be sorted:	
Sort with shellSort, Time consumed:	0.196 (seconds)
Sort with heapSort, Time consumed:	0.220 (seconds)
Sort with mergeSort, Time consumed:	0.156 (seconds)
Sort with quickSort, Time consumed:	0.033 (seconds)

- The input is reversely sorted

3000000 numbers to be sorted:	
Sort with shellSort, Time consumed:	0.306 (seconds)
Sort with heapSort, Time consumed:	0.217 (seconds)
Sort with mergeSort, Time consumed:	0.179 (seconds)
Sort with quickSort, Time consumed:	0.042 (seconds)

Two $\Theta(n)$ sorting algorithms

- Only applicable for **special** cases, but not general cases
- BinSort
- Radix Sort

BinSort Motivation

- Consider $n=5$ integers to be sorted:
 - $A[5]=1, 5, 4, 9, 2$
 - Notice that the **maximum** number is $< 2n = 10$
- Allocate an array $\text{Bin}[10]$
- Place $A[i]$ to $\text{Bin}[A[i]]$, e.g.,
 - Place $A[1] = 5$ to $\text{Bin}[5]$ by setting $\text{Bin}[5] = 1$
 - The other values in Bin are **0**

0	1	1	0	1	1	0	0	0	1	
index	0	1	2	3	4	5	6	7	8	9

BinSort Motivation

- $A[5]=1, 5, 4, 9, 2$

0	1	1	0	1	1	0	0	0	1	
index	0	1	2	3	4	5	6	7	8	9

- BinSort has three steps:
 1. Set $\text{Bin}[j]=0$ for $0 \leq j \leq 9$
 2. Scan array A, set $\text{Bin}[A[i]]=1$ for $0 \leq i \leq 5$
 3. Scan array Bin from the **leftmost to rightmost**, if $\text{Bin}[j] = 1$, number j is in array A, and **output j**
- The **output** is the sequence:1, 2, 4, 5, 9

Binsort

- $A[0, \dots, n-1]$,
- Assume that $A[i] \geq 0$ and $A[i] < c^*n$, c is a constant, e.g., $c = 2$
- Allocate an array B with size c^*n

```
for (j=0; j<c*n; j++)
    B[j]=0;
for (i=0; i<n; i++)
    ++B[ A[i] ]; // may have duplicate numbers
i=0; // the ith sorted number
for (j=0; j<c*n; j++)//number j appears B[j] times
    for (k=0; k<B[j]; k++, i++)
        A[i] = j;
```

- Time complexity
 - $\Theta(cn) + \Theta(n) + \Theta(cn+n)$
 - $= \Theta(cn) = \Theta(n)$, as c is a constant

The application of BinSort is limited

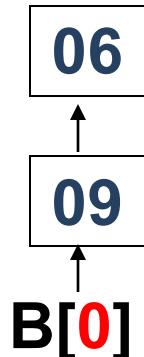
- $A[0, \dots, n-1]$,
- BinSort is applicable when $A[i] < c^*n$
- Consider another example with $n=9$ numbers
 - 09, 85, 68, 86, 47, 06, 39, 34, 30
 - The maximum number 86 is about 10 times larger than n , $\geq n^2=81$
 - If BinSort is applied, an array B with size $87 \geq n^2$ is needed
 - The time complexity then is in $\Omega(n^2)$

Radix Sort -- Extend BinSort

- Some examples of **radix** or **base**
- Radix **10**: the values of each digit may be 0, 1, 2, ..., 9
 - $5_{10}, 16_{10}, 20_{10}, \dots$
- Radix **2**: each digit is 0 or 1
 - $101_2, 10000_2, 10100_2, \dots$
- Radix **26** (26 letters): a, b, c, ..., x, y, z
 - Strings `type', 'alpha', 'go'

09 85 68 86 47 06 39 34 30

out of order

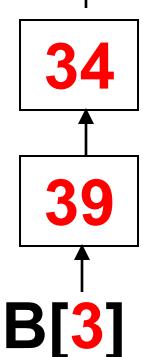


B[0]

B[1]

B[2]

30 out of order



B[3]

B[4]

30 out of order In order

68

86
85

B[5]

B[6]

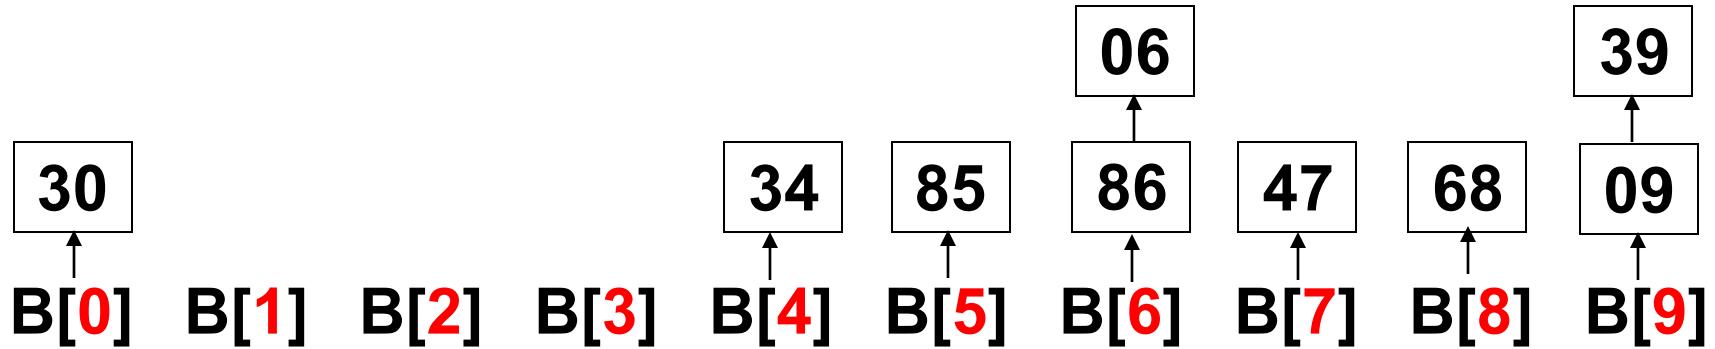
B[7]

B[8]

B[9]

- Each number has two digits
- If we first sort by the **highest** digit:
- But the numbers in the same bin may be **out of order**

09 85 68 86 47 06 39 34 30



- What if we first sort by the lowest digit
- We then collect the numbers in the bins
- The numbers with the **same highest digit** are in **order**, see the numbers with the **same color**

30 34 85 86 06 47 68 09 39

30 34 85 86 06 47 68 09 39

In order

In order

09
06

B[0] B[1]

39
34
30
47

B[3] B[4]

68

B[6]

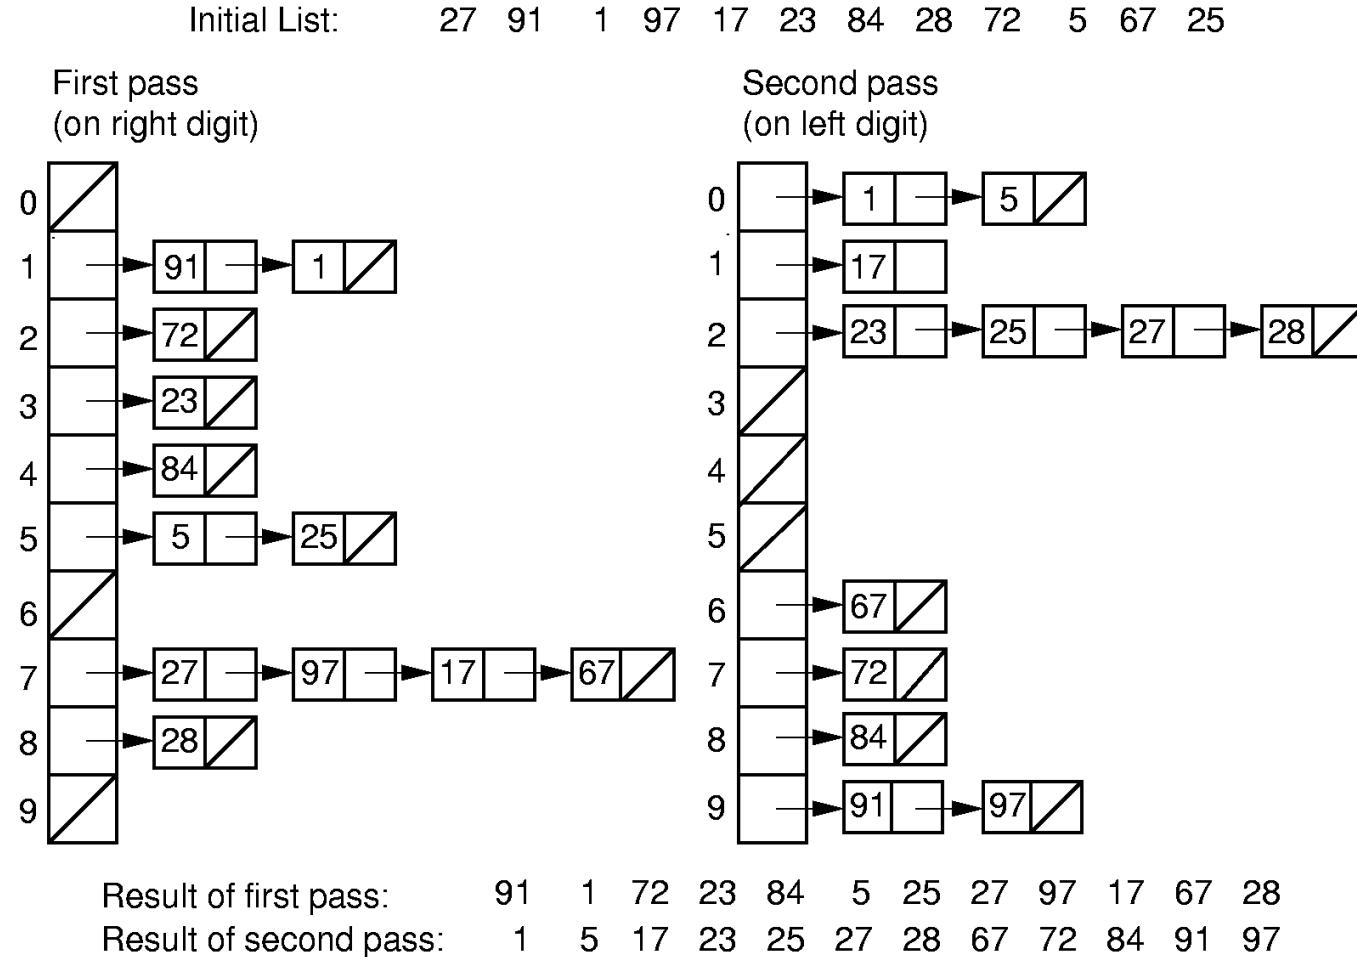
86
85

B[8] B[9]

- We then **sort** the numbers by the **highest** digit
- Collect the numbers in the bins again
- The numbers are **in order** now

06 09 30 34 39 47 68 85 89

RadixSort: sort from the **lowest** digit to the **highest** digit



Radix Sort Cost

- Consider n numbers $A[0, 1, \dots, n-1]$ with radix r , each number has no more than k digits
- Has k BinSorts, from the lowest to the highest
- Each sort takes time $\Theta(n+r)$
- Total Cost: $\Theta(k(n+r))$
- If n numbers are distinct, $k \geq \log_r n$
- If r is small, e.g., $r=2$, radixSort is in $\Theta(n \log n)$
- We usually use large values of r , e.g., $r=1K, 1M$, or even, n

Running time comparisons ($n=3M$)

- random input
- RadixSort is faster if r is larger, $10 \leq r \leq 100k$
- But does not improve any more when r approaches to n

```
3000000 numbers to be sorted:  
Sort with shellSort, Time consumed: 1.062 (seconds)  
Sort with heapSort, Time consumed: 0.953 (seconds)  
Sort with mergeSort, Time consumed: 0.438 (seconds)  
Sort with quickSort, Time consumed: 0.328 (seconds)  
Sort with radixSort (r=10), Time consumed: 0.313 (seconds)  
Sort with radixSort (r=100), Time consumed: 0.156 (seconds)  
Sort with radixSort (r=1000), Time consumed: 0.141 (seconds)  
Sort with radixSort (r=10000), Time consumed: 0.093 (seconds)  
Sort with radixSort (r=100000), Time consumed: 0.078 (seconds)  
Sort with radixSort (r=1000000), Time consumed: 0.079 (seconds)
```

Running time comparisons ($n=3M$)

- The input is already sorted
 - quickSort is faster than radixSort

```
3000000 numbers to be sorted:  
Sort with shellSort, Time consumed: 0.172 (seconds)  
Sort with heapSort, Time consumed: 0.234 (seconds)  
Sort with mergeSort, Time consumed: 0.156 (seconds)  
Sort with quickSort, Time consumed: 0.031 (seconds)  
Sort with radixSort (r=10), Time consumed: 0.329 (seconds)  
Sort with radixSort (r=100), Time consumed: 0.156 (seconds)  
Sort with radixSort (r=1000), Time consumed: 0.156 (seconds)  
Sort with radixSort (r=10000), Time consumed: 0.141 (seconds)  
Sort with radixSort (r=100000), Time consumed: 0.109 (seconds)  
Sort with radixSort (r=1000000), Time consumed: 0.063 (seconds)
```

- The input is reversely sorted

```
3000000 numbers to be sorted:  
Sort with shellSort, Time consumed: 0.281 (seconds)  
Sort with heapSort, Time consumed: 0.234 (seconds)  
Sort with mergeSort, Time consumed: 0.172 (seconds)  
Sort with quickSort, Time consumed: 0.032 (seconds)  
Sort with radixSort (r=10), Time consumed: 0.313 (seconds)  
Sort with radixSort (r=100), Time consumed: 0.156 (seconds)  
Sort with radixSort (r=1000), Time consumed: 0.156 (seconds)  
Sort with radixSort (r=10000), Time consumed: 0.156 (seconds)  
Sort with radixSort (r=100000), Time consumed: 0.110 (seconds)  
Sort with radixSort (r=1000000), Time consumed: 0.078 (seconds)
```

The limitation of RadixSort

- Only applicable to sorting integers
- But **inapplicable** for
 - real numbers
 - Strings has arbitrarily length
 - E.g., **short** string `a', **long** string `dfdfldlfdfdfldjfjlslfjsdfdfdoojll'
 - ...

Lower Bound for Sorting

- We would like to know a lower bound for **all possible** sorting algorithms
- Sorting is $O(n \log n)$ (average, worst cases) because we know algorithms with this upper bound, e.g., **MergeSort** or **HeapSort**
- Sorting takes $\Omega(n)$ time, as each number must be accessed at least once
- Is there any one **better** than $\Theta(n \log n)$?
- It is proved that sorting is $\Omega(n \log n)$
- **MergeSort** and **HeapSort** are **asymptotically optimal** !

Chapter III-8. File Processing and External Sorting

Primary vs. Secondary Storage

- Primary storage: Main memory (RAM)
 - volatile, i.e., data is lost if powered off
 - Usually a few **GB**
 - Expensive (unit: **\$/MB**), fast
- Secondary Storage: Peripheral devices
 - Hard Disk, Solid State Drive (SSD), USB, CD, Tape,...
 - **Non-volatile**
 - Hundreds of **GB**, or **TB**
 - Cheap and slow

Performance Comparisons (typical values)

	Sequential read	seq. write	Random read	Random write
RAM	5 GB/s	4 GB/s	300 MB/s	250 MB/s
Hard Disk	80 MB/s	80 MB/s	0.3 MB/s	0.5 MB/s
SSD	200 MB/s	80 MB/s	25 MB/s	70 MB/s

- Performance of hard disks is **terribly poor** for **random** read and write

Golden Rule of File Processing

- Minimize the number of disk accesses!
 - Arrange information so that you get what you want with few disk accesses
 - Store data on adjacent tracks, rather than randomly
 - Arrange information to minimize future disk accesses
 - Cache

External Sorting

- Problem: Sorting data sets too large to fit into main memory.
 - Assume data are stored on disk drive.
- To sort, portions of the data must be brought into main memory, processed, and returned to disk.
- An external sort should minimize disk accesses.

Model of External Computation

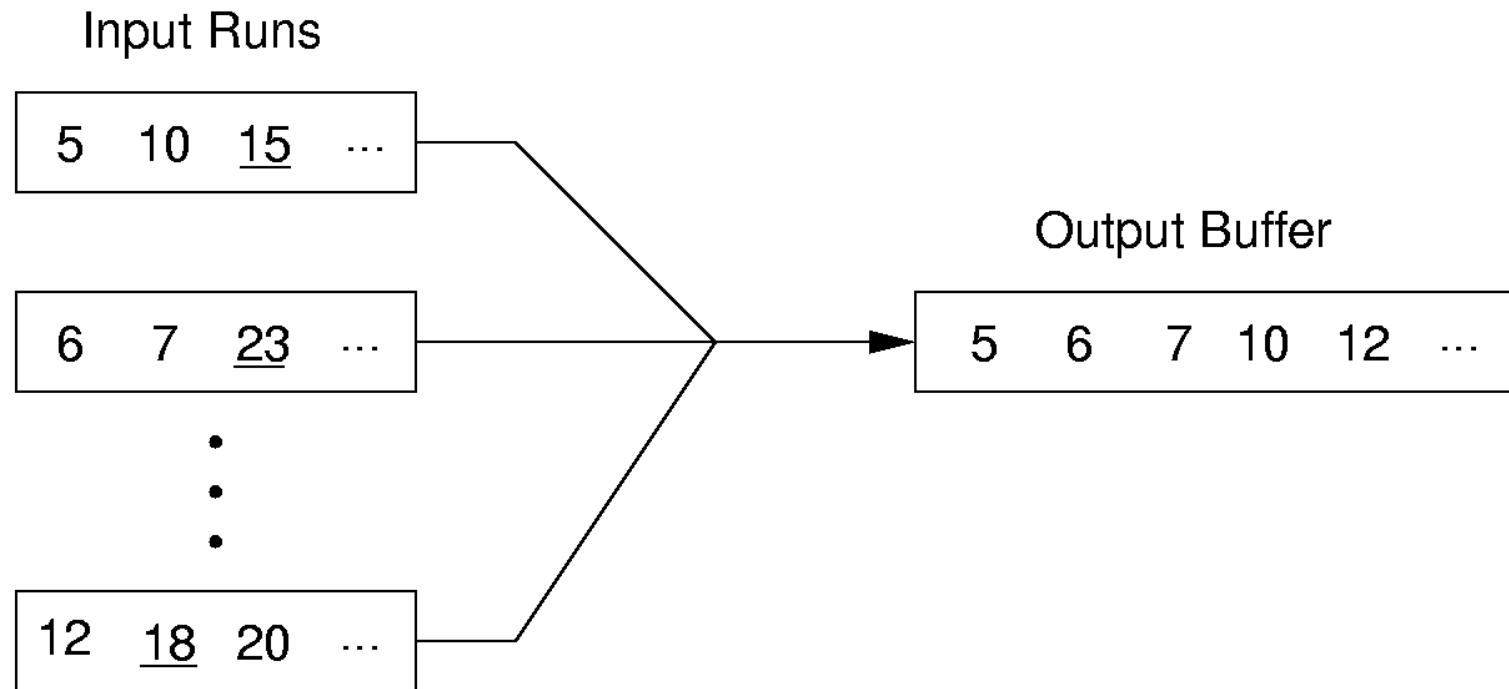
- As **sequential** access is much more efficient than **random** access to the file
 - adjacent logical blocks of the file must be physically adjacent.

External Sorting

- Three Steps:
 1. **Break** a large file into multiple **small** initial blocks, so that each block can be fit into memory
 - E.g., break a 10 GB file into 10 blocks with each being 1 GB
 2. Sorting the blocks by a fast **internal sorting** algorithm one by one, and write back to hard disks
 3. **Merge** the sorted blocks together to form a single sorted file.

Multiway Merge

- Merge **multiple** blocks together, not just two blocks as the internal MergeSort



Summary of Today's Lecture

- Internal sorting
 - THREE basic sorting algorithms ← $\Theta(n^2)$
 - Insertion / bubble / selection sorts
 - ONE medium sorting algorithms ← $\Theta(n^{1.5})$
 - Shell sort
 - THREE fast sorting algorithms ← $\Theta(n \log n)$
 - Merge / quick / heap sorts
 - Two special cases ← $\Theta(n)$
 - Bin / radix sorts
- External sorting

Homework 3

- See course webpage
- **Deadline:** midnight before next lecture
- Submit to: cs_scu@foxmail.com
- File name format:
 - CS311_Hw3_yourID_yourLastName.doc (or .pdf)